CCP Margin Adequacy, Backtesting, and Risk Governance Engine

Detailed LaTeX Explanation with Mathematical Concepts

April 16, 2026

Contents

1 Executive Summary **3**

2 Notation and System Structure **3**
- 2.1 Indices and State Variables . . . 3
- 2.2 Architecture . . . 4

3 Synthetic Data Layer **4**
- 3.1 Risk Factors . . . 4
- 3.2 Realised Volatility Proxy . . . 5
- 3.3 Member Archetypes and Positions . . . 5
- 3.4 Collateral . . . 5

4 Pricing and Daily P&L **5**
- 4.1 Futures Pricing . . . 5
- 4.2 Listed Option Pricing . . . 6
- 4.3 Portfolio Value and Daily P&L . . . 6

5 Scenario Engine and Baseline Risk **6**
- 5.1 Historical Return Matrix . . . 6
- 5.2 Historical Simulation VaR . . . 6
- 5.3 Stressed Window Selection . . . 7
- 5.4 Baseline Margin . . . 7

6 Liquidity Add-On **7**
- 6.1 Motivation . . . 7
- 6.2 Spread Cost . . . 7
- 6.3 Market Impact . . . 8
- 6.4 Liquidation-Horizon Scaling . . . 8
- 6.5 Member Liquidity Multiplier . . . 8

7 Concentration Add-On **8**
- 7.1 ADV Fraction . . . 8
- 7.2 Piecewise Concentration Rate . . . 8

8 Required Margin and Adequacy **9**
8.1 Liquidation-Adjusted Loss . . . 9
8.2 Coverage Ratio and Buffer . . . 9
8.3 Traffic-Light Classification . . . 9
8.4 Margin Call Amount . . . 10

9 Backtesting and Model Monitoring **10**
9.1 Exception Definition . . . 10
9.2 Rolling Exception Count . . . 10
9.3 Interpretation . . . 10

10 Data-Quality and Pricing Controls **11**
10.1 Stale Data . . . 11
10.2 Outlier Returns . . . 11
10.3 Implausible Volatility . . . 11
10.4 Exposure Jumps . . . 11
10.5 Pricing Validity . . . 11

11 Governance and Escalation **11**
11.1 Rule-Based Escalation Framework . . . 11
11.2 Why Governance Matters Mathematically . . . 12

12 Reporting Layer **12**
12.1 Daily Outputs . . . 12
12.2 Weekly and Monthly Outputs . . . 12

13 End-to-End Daily Pipeline **13**

14 Limitations **13**

15 Possible Extensions **14**

16 Appendix: Parameter Summary **14**

17 Appendix: Mapping Equations to Code **15**

1 Executive Summary

The CCP Margin Adequacy, Backtesting, and Risk Governance Engine is a stylised risk platform that evaluates whether a clearing member's posted collateral is sufficient to absorb liquidation-adjusted losses under both ordinary and stressed market conditions. The engine is deliberately broader than a stand-alone Value-at-Risk calculation. It combines:

- synthetic market and portfolio generation,
- full revaluation pricing for futures and listed options,
- historical simulation and stressed Value-at-Risk,
- liquidity and concentration add-ons,
- adequacy controls and margin-call logic,
- backtesting and data-quality controls,
- rule-based governance escalations, and
- committee-ready reporting.

The core business question is:

Is posted collateral sufficient to cover liquidation-adjusted loss for each member on each day?

The engine is educational rather than production-grade, but it is organised in a production-style architecture. The quantitative point is not merely to estimate market risk, but to transform risk estimates into operational and governance decisions.

2 Notation and System Structure

2.1 Indices and State Variables

Let:

$$\begin{aligned} m &\in \{1, \dots, M\} && \text{index clearing members,} \\ i &\in \{1, \dots, I\} && \text{index instruments,} \\ j &\in \{1, \dots, J\} && \text{index risk factors,} \\ t &\in \{1, \dots, T\} && \text{index business dates.} \end{aligned}$$

The key state variables are:

$$\begin{aligned} S_{j,t} &:= \text{spot level of risk factor } j \text{ at date } t, \\ r_{j,t} &:= \log(S_{j,t}/S_{j,t-1}) \text{ one-day log return,} \\ q_{m,i,t} &:= \text{position quantity of member } m \text{ in instrument } i, \\ V_{m,t} &:= \text{portfolio value of member } m \text{ at date } t, \\ \Pi_{m,t} &:= \text{one-day profit and loss of member } m, \\ \mathrm{IM}_{m,t} &:= \text{required initial margin,} \\ \mathrm{PM}_{m,t} &:= \text{posted collateral after haircut.} \end{aligned}$$

2.2 Architecture

The implementation is separated into the following modules.

Module	Function	Main Files
Data and portfolio	Generate synthetic market data, instruments, member types, positions, collateral	`src/portfolio.py`, `src/data_loader.py`
Pricing and P&L	Full revaluation of futures and options; daily P&L generation	`src/pricing.py`
Scenario engine	Historical and stressed return scenario construction	`src/scenarios.py`
Margin methodology	Historical VaR, stressed VaR, liquidity and concentration add-ons	`src/margin.py`, `src/liquidity.py`, `src/concentration.py`
Adequacy controls	Coverage ratios, traffic lights, margin call logic	`src/controls.py`
Monitoring	Backtesting exceptions and data-quality checks	`src/backtesting.py`
Governance	Escalation rules and breach handling	`src/escalation.py`
Reporting	Daily, weekly, and committee-style outputs	`src/reporting.py`

3 Synthetic Data Layer

3.1 Risk Factors

The platform models eight primary risk factors:

$$\{\mathrm{SPX}, \mathrm{NDX}, \mathrm{RTY}, \mathrm{TY}, \mathrm{FV}, \mathrm{US}, \mathrm{VIX}, \mathrm{CL}\}.$$

These cover equity-index, rates, volatility, and commodity risk. The synthetic market-data generator uses fat-tailed shocks rather than Gaussian shocks. Specifically, each factor is generated using Student-t innovations with $\nu = 5$ degrees of freedom:

$$\varepsilon_{j,t} \sim t_\nu, \qquad \nu = 5.$$

If σ_j^{ann} denotes the annual volatility assigned to factor j, the daily-scale volatility is:

$$\sigma_j^{\mathrm{day}} = \frac{\sigma_j^{\mathrm{ann}}}{\sqrt{252}}.$$

The synthetic one-day return is then:

$$r_{j,t} = \sigma_j^{\mathrm{day}} \varepsilon_{j,t},$$

and the spot path is constructed through exponential accumulation:

$$S_{j,t} = S_{j,0} \exp\left(\sum_{u=1}^{t} r_{j,u}\right).$$

This is useful because the Student-t innovations generate heavier tails than a Gaussian model, which makes downstream VaR and backtesting behaviour more realistic.

3.2 Realised Volatility Proxy

The option-pricing layer uses a realised-volatility proxy. For each factor:

$$\widehat{\sigma}_{j,t}^{(20)} = \sqrt{252}\,\mathrm{Std}\left(r_{j,t-19}, \ldots, r_{j,t}\right),$$

where the rolling standard deviation is computed over a 20-business-day lookback. This realised volatility is used as the volatility input for Black-Scholes in the listed-option module.

3.3 Member Archetypes and Positions

The engine uses ten stylised members, including directional macro, relative value, concentrated volatility, diversified, and weak-liquidity members. The design purpose is not only to create heterogeneity of exposures, but also to ensure that the governance layer has something non-trivial to detect.

At each date t, a member portfolio is represented by quantities $q_{m,i,t}$. Market value for each position is:

$$\mathrm{MV}_{m,i,t} = q_{m,i,t} \cdot M_i \cdot S_{\mathrm{und}(i),t},$$

where M_i is the contract multiplier and $\mathrm{und}(i)$ is the underlying risk factor for instrument i.

3.4 Collateral

Each member posts a mixture of cash and government bond collateral. Let:

$$\begin{aligned} C_{m,t}^{\mathrm{cash}} &:= \text{cash collateral}, \\ C_{m,t}^{\mathrm{bond}} &:= \text{government bond collateral}, \\ h_{m,t} &:= \text{bond haircut}. \end{aligned}$$

Then posted margin after haircut is:

$$\mathrm{PM}_{m,t} = C_{m,t}^{\mathrm{cash}} + C_{m,t}^{\mathrm{bond}}(1 - h_{m,t}).$$

This quantity is what the adequacy layer compares against required initial margin.

4 Pricing and Daily P&L

4.1 Futures Pricing

For a futures contract i on underlying j, the engine uses a simple linear mark:

$$P_{i,t}^{\mathrm{future}} = M_i S_{j,t}.$$

For quantity $q_{m,i,t}$, the contribution to portfolio value is:

$$V_{m,i,t}^{\mathrm{future}} = q_{m,i,t} P_{i,t}^{\mathrm{future}}.$$

4.2 Listed Option Pricing

For listed options, the engine uses the Black-Scholes model. Let S be underlying spot, K the strike, T time to expiry in years, r the risk-free rate, and σ the volatility input. Define:

$$d_1 = \frac{\ln(S/K) + \left(r + \frac{1}{2}\sigma^2\right) T}{\sigma\sqrt{T}}, \qquad d_2 = d_1 - \sigma\sqrt{T}.$$

Then the call and put prices are:

$$\begin{aligned} C(S, K, T, r, \sigma) &= S\Phi(d_1) - Ke^{-rT}\Phi(d_2), \\ P(S, K, T, r, \sigma) &= Ke^{-rT}\Phi(-d_2) - S\Phi(-d_1). \end{aligned}$$

The marked option value in the engine is:

$$P_{i,t}^{\text{option}} = M_i \times \begin{cases} C(S, K, T, r, \sigma), & \text{if } i \text{ is a call,} \\ P(S, K, T, r, \sigma), & \text{if } i \text{ is a put.} \end{cases}$$

4.3 Portfolio Value and Daily P&L

Portfolio value is the sum over all open instruments:

$$V_{m,t} = \sum_{i=1}^{I} q_{m,i,t} P_{i,t}.$$

The one-day P&L used throughout the engine is the full-revaluation difference:

$$\Pi_{m,t} = V_{m,t} - V_{m,t-1}.$$

This is important: the VaR is not built from a linear factor approximation at the member level, but from the realised full-revaluation P&L series generated by the pricing layer.

5 Scenario Engine and Baseline Risk

5.1 Historical Return Matrix

Let the daily return vector be:

$$R_t = (r_{1,t}, r_{2,t}, \ldots, r_{J,t})^\top.$$

The historical scenario engine constructs the matrix of the trailing 500 business days:

$$\mathcal{H}_t = \{R_{t-499}, R_{t-498}, \ldots, R_t\}.$$

This forms the basis for historical simulation VaR.

5.2 Historical Simulation VaR

For a fixed member m and date t, define the trailing member-level P&L sample:

$$\mathcal{P}_{m,t} = \{\Pi_{m,t-499}, \ldots, \Pi_{m,t}\}.$$

The 99% Historical Simulation VaR is:

$$\text{HSVaR}_{m,t}^{99\%} = -Q_{0.01}\left(\mathcal{P}_{m,t}\right),$$

where $Q_{0.01}$ is the empirical 1% quantile. Since losses correspond to negative P&L, the minus sign converts the lower-tail quantile into a positive capital number.

5.3 Stressed Window Selection

The engine also computes a stressed VaR using a 60-day stressed sub-window. The stressed period is selected as the contiguous window with the highest realised SPX volatility. If $r_{\mathrm{SPX},u}$ denotes SPX returns, the stressed-window endpoint can be expressed conceptually as:

$$u^\star = \arg\max_u \mathrm{Std}\left(r_{\mathrm{SPX},u-59}, \ldots, r_{\mathrm{SPX},u}\right).$$

The stressed set of dates is then:

$$\mathcal{S}_t = \{u^\star - 59, \ldots, u^\star\}.$$

The stressed VaR is computed by restricting the member P&L series to this stressed subset:

$$\mathrm{SVaR}_{m,t}^{99\%} = -Q_{0.01}\left(\{\Pi_{m,u} : u \in \mathcal{S}_t\}\right).$$

If too few stressed observations overlap the available member P&L history, the implementation falls back to the ordinary historical VaR.

5.4 Baseline Margin

The baseline margin is:

$$B_{m,t} = \max\left(\mathrm{HSVaR}_{m,t}^{99\%}, \mathrm{SVaR}_{m,t}^{99\%}\right).$$

This is a simple but important anti-procyclicality device. During calm periods, ordinary historical VaR may fall materially. The stressed VaR acts as a floor tied to a high-volatility regime.

6 Liquidity Add-On

6.1 Motivation

Pure VaR measures mark-to-market loss but not the cost of actually closing out a portfolio. A CCP margin framework should reflect liquidation cost. The engine therefore augments the baseline margin with a liquidity add-on that contains:

- half-spread crossing cost,
- market impact, and
- liquidation-horizon scaling.

6.2 Spread Cost

For position $q_{m,i,t}$, underlying spot $S_{j,t}$, bid-ask spread in basis points $b_{j,t}$, and multiplier M_i, the spread-crossing component is:

$$\mathrm{SC}_{m,i,t} = |q_{m,i,t}|\, M_i\, S_{j,t}\, \frac{b_{j,t}}{20000}.$$

The denominator 20000 arises because the engine uses half the full bid-ask spread, and basis points must be converted into decimals.

6.3 Market Impact

Let $\mathrm{ADV}_{j,t}$ be the average daily volume of the relevant underlying, and let α be the impact coefficient. The market-impact estimate is:

$$\mathrm{MI}_{m,i,t} = \alpha \sqrt{\frac{|q_{m,i,t}|}{\mathrm{ADV}_{j,t}}}\, S_{j,t}\, |q_{m,i,t}|\, M_i.$$

In the implementation, $\alpha = 0.10$.

The square-root form follows the common empirical intuition that execution cost increases sublinearly with size at first, but materially once trade size becomes large relative to market depth.

6.4 Liquidation-Horizon Scaling

If ℓ_j is the liquidation horizon in days for factor j, then the raw liquidity cost is scaled by $\sqrt{\ell_j}$:

$$\mathrm{LC}_{m,i,t} = (\mathrm{SC}_{m,i,t} + \mathrm{MI}_{m,i,t})\sqrt{\ell_j}.$$

The engine uses:

$$\ell_j = \begin{cases} 2, & \text{liquid bucket,} \\ 3, & \text{semi-liquid bucket,} \\ 5, & \text{illiquid bucket.} \end{cases}$$

6.5 Member Liquidity Multiplier

Each member type has an additional liquidity multiplier λ_m^{liq} to reflect expected liquidation difficulty:

$$\lambda_m^{\mathrm{liq}} \in \{0.7, 0.8, 1.0, 1.3, 1.5\},$$

with larger values assigned to concentrated-volatility or weak-liquidity members.

The total liquidity add-on is:

$$\mathrm{LA}_{m,t} = \lambda_m^{\mathrm{liq}} \sum_{i=1}^{I} \mathrm{LC}_{m,i,t}.$$

7 Concentration Add-On

7.1 ADV Fraction

Concentration is measured using position size relative to average daily volume:

$$f_{m,i,t} = \frac{|q_{m,i,t}|}{\mathrm{ADV}_{j,t}}.$$

7.2 Piecewise Concentration Rate

The engine defines a concentration rate function:

$$c(f) = \begin{cases} 0, & f < 0.10, \\ 0.10, & 0.10 \le f < 0.20, \\ 0.25, & f \ge 0.20. \end{cases}$$

For each member-day, the implementation takes the maximum concentration rate across instruments:

$$c_{m,t}^{\max} = \max_i c(f_{m,i,t}).$$

The concentration add-on is then:

$$\mathrm{CA}_{m,t} = c_{m,t}^{\max} \cdot B_{m,t}.$$

This design means concentration does not stack instrument by instrument; rather, the most severe crowding signal determines the add-on rate applied to baseline margin.

8 Required Margin and Adequacy

8.1 Liquidation-Adjusted Loss

The central quantity in the project is the liquidation-adjusted loss:

$$\mathrm{LAL}_{m,t} = B_{m,t} + \mathrm{LA}_{m,t} + \mathrm{CA}_{m,t}.$$

The engine sets:

$$\mathrm{IM}_{m,t} = \mathrm{LAL}_{m,t}.$$

Equivalently:

$$\mathrm{IM}_{m,t} = \max\left(\mathrm{HSVaR}_{m,t}^{99\%}, \mathrm{SVaR}_{m,t}^{99\%}\right) + \mathrm{LA}_{m,t} + \mathrm{CA}_{m,t}.$$

8.2 Coverage Ratio and Buffer

Once posted collateral is known, adequacy is measured through:

$$\mathrm{Coverage}_{m,t} = \frac{\mathrm{PM}_{m,t}}{\mathrm{IM}_{m,t}}, \qquad \mathrm{Buffer}_{m,t} = \mathrm{PM}_{m,t} - \mathrm{IM}_{m,t}.$$

8.3 Traffic-Light Classification

The traffic-light rule is:

$$\mathrm{TL}_{m,t} = \begin{cases} \text{green}, & \mathrm{Coverage}_{m,t} \geq 1.10, \\ \text{amber}, & 1.00 \leq \mathrm{Coverage}_{m,t} < 1.10, \\ \text{red}, & \mathrm{Coverage}_{m,t} < 1.00. \end{cases}$$

This creates a direct operational interpretation:

- green means comfortable over-collateralisation,
- amber means coverage is technically sufficient but thin,
- red means posted collateral is insufficient.

8.4 Margin Call Amount

Let the shortfall be:

$$\Delta_{m,t} = \mathrm{IM}_{m,t} - \mathrm{PM}_{m,t}.$$

Let τ be the operational threshold and μ the minimum transfer amount. In the implementation:

$$\tau = 100{,}000, \qquad \mu = 100{,}000.$$

The margin call function is:

$$\mathrm{MC}_{m,t} = \begin{cases} 0, & \Delta_{m,t} \le \tau, \\ \max(\Delta_{m,t}, \mu), & \Delta_{m,t} > \tau. \end{cases}$$

This prevents operational noise from generating very small calls.

9 Backtesting and Model Monitoring

9.1 Exception Definition

The actual realised loss on date t is:

$$\mathrm{AL}_{m,t} = \max(-\Pi_{m,t}, 0).$$

Backtesting compares this realised loss against prior-day required margin. The exception indicator is:

$$I_{m,t} = \mathbf{1}\left\{\mathrm{AL}_{m,t} > \mathrm{IM}_{m,t-1}\right\}.$$

An exception therefore means the prior day's margin estimate was not sufficient to cover the next day's realised loss.

9.2 Rolling Exception Count

For a 250-day rolling window, the engine computes:

$$E_{m,t}^{(250)} = \sum_{u=t-249}^{t} I_{m,u}.$$

The exception traffic light is:

$$\mathrm{BT}_{m,t} = \begin{cases} \text{green}, & E_{m,t}^{(250)} < 3, \\ \text{amber}, & 3 \le E_{m,t}^{(250)} < 5, \\ \text{red}, & E_{m,t}^{(250)} \ge 5. \end{cases}$$

9.3 Interpretation

The monitoring logic is not purely statistical. If exceptions begin to cluster, this suggests a possible regime-specific weakness in the model, for example:

- the VaR window may react too slowly,
- stressed-window selection may be insufficiently conservative,
- liquidity or concentration add-ons may be too weak, or
- the options pricing proxy may fail to capture volatility shocks.

10 Data-Quality and Pricing Controls

10.1 Stale Data

For each factor, the engine flags sequences of zero returns. If:

$$|r_{j,t}| < \varepsilon$$

for more than two consecutive business days, a stale-data flag is raised. Conceptually, this is a control against unchanged vendor marks rather than a market event.

10.2 Outlier Returns

An outlier-return flag is triggered if:

$$|r_{j,t}| > 0.15.$$

10.3 Implausible Volatility

The realised-volatility input is flagged if:

$$\widehat{\sigma}_{j,t}^{(20)} < 0.01 \qquad \text{or} \qquad \widehat{\sigma}_{j,t}^{(20)} > 3.00.$$

10.4 Exposure Jumps

For member-instrument market values, the engine computes day-over-day changes. If a new change exceeds the historical mean plus three standard deviations,

$$|\Delta\mathrm{MV}_{m,i,t}| > \mu_{m,i}^{\Delta\mathrm{MV}} + 3\sigma_{m,i}^{\Delta\mathrm{MV}},$$

an exposure-jump flag is raised.

10.5 Pricing Validity

Option prices must remain finite and non-negative:

$$P_{i,t}^{\text{option}} \in [0, \infty), \qquad P_{i,t}^{\text{option}} \neq \mathrm{NaN}, \qquad P_{i,t}^{\text{option}} \neq \infty.$$

These controls do not change the economics of the model, but they are essential for a governed calculation pipeline.

11 Governance and Escalation

11.1 Rule-Based Escalation Framework

The governance layer turns quantitative states into actions. Define indicator functions:

$$\begin{aligned}
R_{m,t}^{\text{red}} &= \mathbf{1}\{\mathrm{TL}_{m,t} = \text{red}\}, \\
R_{m,t}^{\text{call}} &= \mathbf{1}\{\mathrm{MC}_{m,t} > 0\}, \\
R_{m,t}^{\text{bt}} &= \mathbf{1}\{E_{m,t}^{(250)} \geq 4\}, \\
R_{m,t}^{\text{conc}} &= \mathbf{1}\left\{\frac{\mathrm{CA}_{m,t}}{\mathrm{IM}_{m,t}} > 0.25\right\}, \\
R_{t}^{\text{stale}} &= \mathbf{1}\{\text{critical stale-data condition holds at } t\}.
\end{aligned}$$

The engine's main escalation rules can then be summarised as:

Rule	Quantitative condition	Owner
Single red breach	$R^{\text{red}}_{m,t} = 1$	Risk Analyst
Consecutive red days	two or more consecutive dates with $R^{\text{red}}_{m,t} = 1$	Senior Risk Officer
Methodology review	$E^{(250)}_{m,t} \geq 4$	Model Validation
Committee watchlist	$\text{CA}_{m,t}/\text{IM}_{m,t} > 0.25$	Risk Committee
Provisional run	$R^{\text{stale}}_{t} = 1$	Market Data Team
Operational margin call	$R^{\text{call}}_{m,t} = 1$	Margin Operations

11.2 Why Governance Matters Mathematically

Mathematically, the model is a function:

$$\mathcal{M} : (\text{market data}, \text{positions}, \text{collateral}, \text{parameters}) \mapsto (\text{margin}, \text{adequacy}, \text{exceptions}, \text{escalations}).$$

The governance framework defines how to act when the output of $\mathcal{M}$ breaches certain thresholds. In other words, governance is the decision layer attached to the model.

12 Reporting Layer

The reporting module transforms the state of the engine into artifacts intended for daily and committee consumption.

12.1 Daily Outputs

For the latest date t, the daily member-adequacy table contains:

$$(\text{IM}_{m,t}, \text{PM}_{m,t}, \text{Coverage}_{m,t}, \text{TL}_{m,t}, \text{MC}_{m,t}).$$

The breach register supplements this with:

$$(\text{breach type}, \text{severity}, \text{owner}, \text{escalation level}, \text{target resolution date}).$$

12.2 Weekly and Monthly Outputs

The weekly report aggregates:

- backtesting exceptions,
- concentration events,
- stale-data incidents, and
- escalation counts.

The monthly committee pack adds:

- time series of total required margin,

- weakest members by average coverage,
- coverage-ratio distribution statistics,
- exception summaries, and
- known limitations and recommended actions.

13 End-to-End Daily Pipeline

At a high level, the daily pipeline can be described as:

1. Generate or load market data, instruments, members, positions, and collateral.
2. Revalue all portfolios and compute member-level daily P&L.
3. Compute historical and stressed VaR for each member.
4. Add liquidity and concentration penalties.
5. Compute required margin and compare it with posted collateral.
6. Generate traffic lights and margin calls.
7. Backtest prior-day margins against realised losses.
8. Run data-quality and model-control checks.
9. Generate escalation events.
10. Produce daily, weekly, and committee-facing reports.

This can be summarised compactly as:

$$\begin{aligned}&(\text{Data}) \to (\text{Pricing}) \to (\text{VaR}) \to (\text{Add-ons})\\ &\qquad \to (\text{Adequacy}) \to (\text{Backtesting and Controls}) \to (\text{Escalation}) \to (\text{Reporting}).\end{aligned}$$

14 Limitations

The mathematical structure is coherent, but several simplifications remain:

- The margin methodology is stylised and omits many CCP production features such as cross-margin offsets, legal account segregation, and default waterfalls.
- Option pricing uses a Black-Scholes framework with realised-volatility proxy rather than a full implied-volatility surface.
- Rates futures are handled through a linear approximation rather than a richer term-structure model.
- Liquidity and concentration costs are intentionally simple and parameterised rather than calibrated from execution data.
- Default management and auction mechanics are not modelled.

These limitations do not invalidate the project. Rather, they define its role as a teaching and portfolio platform for quantitative risk concepts.

15 Possible Extensions

Natural quantitative extensions include:

- filtered historical simulation or expected shortfall in place of pure historical VaR,
- scenario-dependent implied-volatility surfaces for listed options,
- more realistic liquidation curves and impact calibration,
- wrong-way collateral risk and eligibility schedules,
- member credit-state modelling, and
- richer stress scenarios with cross-asset contagion.

16 Appendix: Parameter Summary

Parameter	Value	Role
VaR confidence level	99%	Tail quantile for margin sizing
Historical window	500 days	Historical simulation sample
Stressed window	60 days	High-volatility sub-window
Risk-free rate	4.5%	Black-Scholes input
Impact coefficient α	0.10	Market-impact scaling
Green adequacy threshold	1.10	Comfortable collateral coverage
Amber adequacy threshold	1.00	Minimum full coverage
Margin threshold	$100,000	Operational trigger threshold
Minimum transfer amount	$100,000	Minimum call size
Backtesting warning level	3	Amber exception level
Backtesting critical level	5	Red exception level
Stale-data threshold	2 days	Data-quality control threshold
Outlier return threshold	15%	Return anomaly control threshold

17 Appendix: Mapping Equations to Code

Concept	Implementation location
Synthetic market generation	`src/portfolio.py::generate_market_data`
Instrument and member generation	`src/portfolio.py::generate_instruments`, `generate_member_profiles`
Full revaluation P&L	`src/pricing.py::compute_daily_pnl`
Black-Scholes pricing	`src/pricing.py::_bs_price`, `price_instrument`
Historical and stressed scenarios	`src/scenarios.py`
Historical VaR and stressed VaR	`src/margin.py::historical_simulation_var`, `stressed_var`
Liquidity add-on	`src/liquidity.py::compute_liquidity_addon`
Concentration add-on	`src/concentration.py::compute_concentration_addon`
Coverage ratio and margin call	`src/controls.py`
Backtesting exceptions	`src/backtesting.py::compute_exceptions`
Escalation rules	`src/escalation.py`
Reports and breach register	`src/reporting.py`